

000
001
002
003
004
005
006
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053054
055
056
057
058
059
060
061
062
063
064
065
066
067
068
069
070
071
072
073
074
075
076
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092
093
094
095
096
097
098
099
100
101
102
103
104
105
106
107

Property Tunable Image Manipulation with Self-Generated Supervision

Anonymous WACV submission

Paper ID 697

Abstract

We propose instance-invariant image synthesis steered by continuously tunable properties. The key challenge is that training data distribution is often sparse, including the extreme case of having only one instance sample. We introduce an infinite synthetically-generated scheme to continuously manipulate the input image. For each training image, the model self-generates its manipulated image with the synthesized desired property label as the self-supervision for the model's training. A key objective is to produce desired property change while maintaining instance-invariance. We implement this using a TunerGAN: infusing parallel streams of input-output messages to the GAN network. The messages, once decoded into feature maps represent tunable properties and instance invariance. The output messages extracted from the synthesized image, are self-checked with the input messages to ensure instance-level synthesis is done correctly: e.g., face rotated correctly with identity preserved. Our framework is general for image manipulation and we have conducted experiments on a variety of tasks to demonstrate its effectiveness.

1. Introduction

Image generation via domain style transfer has demonstrated the ability to create images with desired properties at a category level [38, 1]. One can generalize this approach to multiple-property manipulations by treating property states as different domains/styles, as described in [26, 18, 4]. A less explored but equally important direction is to manipulate specific properties of images at *instance-level*. For example, we want to rotate a person's head without changing his/her identity and expression. This is not only helpful for more controllable image editing but also beneficial to many image understanding tasks such as state-estimation and recognition (identification, categorization, etc.). Learning to manipulate a specific property (e.g. head pose) or multiple parallel properties can be realized by learning two disentangled representations, one for the related properties themselves and one for the other information, which is in-

dependent of these properties. The former representation is facilitated to estimate the state of the selected property (head pose), while the latter is more robust for recognition tasks (e.g. face recognition) due to its invariance to the properties' changes. This is conceptually consistent with the general strategy of perception by synthesizing.

The distinction between category-level vs instance-level image manipulation is important. For category-level image manipulation, the learned model is not required to separate the tuning property from instance invariant information. For example, to change one's facial expression, we would not want to add sunglasses to his/her face. Approaches such as InfoGAN [1] can generate an image with desired property change but often confound related factors together. Three fundamental challenges for instance-level properties tuning are 1) succinct controllability, 2) instance-level descriptiveness, and 3) instance-level discriminativeness.

Succinct controllability means we want to have simple control variables, such as degree of head pose/smile or exact age, to control image generation. MUNIT-like approaches [12] use latent vectors to control image generations would need to automate the generation of the latent vectors to work in this case. This is not an easy task since we don't have data for most of the image manipulation targets. In our algorithm, we self-generate data and learn from it at the same time. Furthermore, we ensure succinct controllability by learning a decoder that turns a tunable property (e.g. pose) into a latent feature vector.

Instance-level descriptiveness means that we want to preserve as many details about that instance (e.g. a person's identity and appearance when tuning his/her head pose). Instance-level discriminativeness means we have a precise checking on the generated images to ensure the desired properties variation while maintaining instance-invariance. The two objectives of descriptiveness and discriminativeness are often in conflict with each other. The richer descriptiveness of the instance details, the harder it is for the discriminativeness check. One way to prevent cheating is to ensure self-generation consistency: check if the descriptiveness feature can generate the original image. The approach in MUNIT [12] seems to offer a solution: it ensures the im-

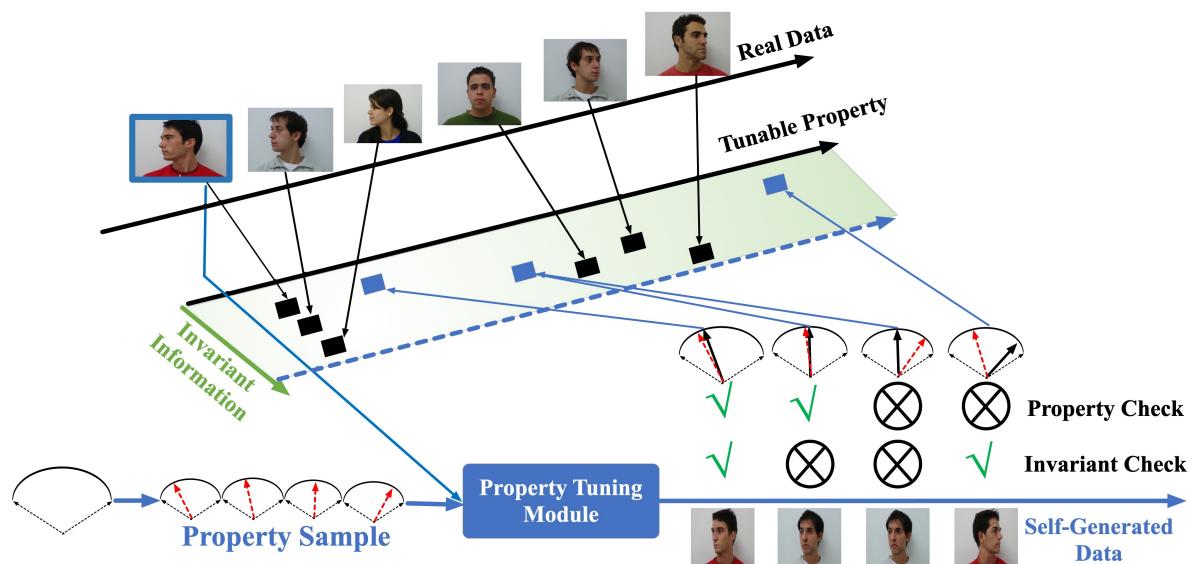
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215

Figure 1. Tunable image manipulation with self-generated supervision. The real data distribution on the top is sparse and unbalanced. For an input image, we disentangle its representation into the tunable property and the invariant information. Then we generate images having different property values by the property tuning module. The disentanglement is checked by whether the generated image is the same person as the source image and whether the property is the same as the sampled property. To show more effectively, we only show the framework to tune one specific property. For tuning multiple properties, the tunable capsule network has multiple parallel outputs, which are responsible for different properties.

ages before-and-after manipulation share the same ‘content’ feature code. However, sharing exactly the same code throw away too many informative features. For example, the nose has unique features different from the frontal to profile view, and yet we want to preserve two different nose details. We solve this using Capsule Network as a discriminativeness checking mechanism. The many-to-one mapping property of the Capsule network provides us the flexibility of having a rich multimodal descriptiveness instance encoding while ensuring precise instance-level discriminativeness.

We formulate instance-level image properties tuning as a message-passing image generator. The message-passing framework explicitly exposes the input streams of 1) succinct controllability, 2) instance-level descriptiveness features, and output stream of 3) instance-level discriminativeness features. These messages play the roles of ‘style’ and ‘content’ code of the MUNIT, but they are succinct and continuously tunable. Inside of the generator, these messages are expanded into expressive multimodal instance-level feature vectors. To realize the above designs, we develop a novel generative adversarial network named TunerGAN, which contains the following three distinctive novelties:

- **Self-Generated Supervision:** the supervision for tuning (pairs of a synthesized label and corresponding manipulated data) is self-generated, as shown in Figure 1. Such a self-supervision frees the model from relying on paired training samples so that any infinite feasible state values can be sampled for training. This also al-

lows training with diverse and arbitrarily sparsely distributed data in the property space (including the extreme case of having only one single sample for some training instances).

- **Capsule Structure for Representation:** two capsule blocks are introduced for learning representations for property-invariant information and property-state information, respectively. Using capsule networks has two advantages: a) its self-attention mechanism allows more robust representation; b) its vector representation can model multimodal distributions to regress complex continuous signals.

- **Multi-Branch Estimator:** property state estimation is critical for effective tuning as it provides feedback for loss checking, and with a multi-branch architecture the state estimator can focus on different aspects and then fuse the results, leading to superior performance.

2. Related Work

2.1. Image Translation and Image Manipulation

In general, image-to-image translation describes a task converting an image with source representation to an image with target representation. Many typical computer vision topics can be summarized as image-to-image translation tasks, including semantic segmentation [21, 33], image restoration and enhancing [22, 35], image editing and

inpainting [6, 25, 32], super resolution [16, 3, 28]. Due to the success of extensions on conditional GANs [24, 13, 30, 4], many works (such as CycleGAN [38], etc. [34, 17, 12]) develop translation models between two distinctive domains, and some others like StarGAN [4] and STGAN [18] can even handle the translation among multiple domains. DRIT [17] and MUNIT [12] introduced a content-preserving loss to achieve representation disentanglement by using two encoders to encode content and style separately, which allows easier style translation. For image manipulation, a representative work [2] first explores multiple paths to perform the translation and generate the manipulated intermediate images between two given samples. However, these methods are limited to the translation among multiple fixed domains or interpolation between two specific samples. In contrast, our TunerGAN can synthesize a manipulated output with targeted property/properties to be of arbitrary value(s), not restricted to those of the actual samples. By tuning property/properties with direct input values whilst preserving the instance-invariant information, TunerGAN can supervise itself via property-specified synthesizing and invariant property value checking after translation, which is not possible in any existing work or their possible combinations (e.g. StarGAN + MUNIT).

2.2. Capsule Network

The concept of capsules is invented in [9] and used recently in [27, 10]. CapsNet [27] is designed for image feature extraction, which is developed based on CNN. Unlike traditional CNN, in which the presence of feature is represented with scalar values in feature maps, the features in CapsNet are represented with capsules (vectors). In [27], the direction of capsules reflects the properties of the features and the length (2-norm) of capsules reflects the probability of the presence of different features. The transmission of information between layers follows the dynamic routing mechanism, which can help us obtain richer information than before. For example, the elements of the capsule may be used for learning different local or global properties, such as color, shape, size, or pose of objects. And the whole length of the capsule can indicate the probability of existence. We adopt capsule architecture to obtain a better global and local representation of visual images.

3. The Proposed Method

This section presents the overall framework and key components, with design motivations and insights explained. Due to the space limits, detailed network structures and implementations are left to the supplementary material.

3.1. Image Manipulation by Property Tuning

Image manipulation can be formulated as domain/style transfer or synthesis tasks. We aim to modulate a tunable

property α (e.g. head pose) on an instance image I_s from one source state α_s to another state α_t , while keeping other properties β unchanged. This can be realized by learning to disentangle two kinds of representations, one for the tunable property α itself and the other for the rest information β which is independent of α . We generate the disentangled representations using a property decoder P , an invariant encoder E and an invariant capsule module C . P projects a low dimensional property value to the property representation so that we can tune the specific property by controlling the value of α . E encodes the invariant properties. The image manipulation with the disentangled feature can be realized by:

$$I_t \equiv G(E(I_s) \pm P(\alpha_t)), \quad (1)$$

where G is the image generation function, I_s is the source image, and I_t is the tuned image.

We propose a self-guided generation of data with the continuous-valued tunable property. It is a significant step because we can now sample an infinite number of labels as well as manipulate any training image according to the self-generated labels. By separating the image information from the tunable properties, we also obtain control over the invariant properties of each sampled image. As shown in Figure 1 and Figure 2, we randomly sample one target state value α_t to generate I_t from the image I_s . Then we optimize an estimator η to measure the state representation for both the real and generated images. To be noted, the two disentangled features are both checked: 1) we check whether the invariant features are the same, *i.e.* $C(E(I_s)) \stackrel{?}{=} C(E(I_t))$, to ensure the generated sample images retain the same identity, where C is the invariant capsule to project the feature representation to low-dimensional capsule output; 2) for the tunable property, we also check its state by $\eta(I_t) \stackrel{?}{=} \alpha_t$. Moreover, we compute the cycle-consistency loss between I_s and re-manipulated image $\hat{I}_s = G(E(I_t) + P(\alpha_s))$.

Training with self-generated data makes it possible to continuously tune the property and efficiently explore the property space. Checking both the invariant and the tunable property allows the representation to be disentangled as expected, so we can manipulate the tunable property while maintaining other properties. An adversarial discriminator is also utilized to teach the generator to focus on the details which are irrelevant to the tunable property, and we adopt a multi-branch structure to stabilize the training.

3.2. Training with Self-generated Data

The training sample for the image manipulation is paired with the corresponding label, where the input image I_s checks with the labeled property α_s by $\eta(I_s) \stackrel{?}{=} \alpha_s$, and the generated image I_t checks with α_t by $\eta(I_t) \stackrel{?}{=} \alpha_t$. As a result, the property label α_s is paired with the train

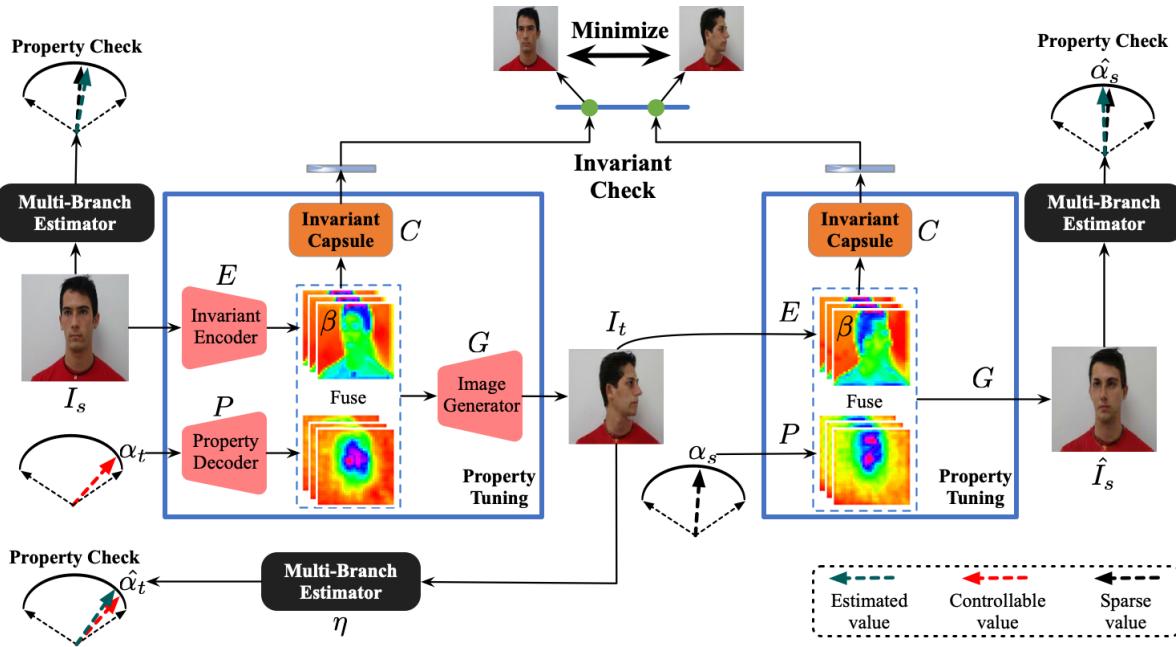


Figure 2. Training the property tuning module with the invariant check and the property check. The image feature from the invariant encoder is converted into the capsule vector by the invariant capsule module. We conduct the invariant checking by minimizing the distance between the capsule vectors. The source and generated images have their properties checked by a multi-branch estimator to get aligned with the labeled or desired properties. Besides the invariant capsule (C), other components of the second reconstruction part including the invariant encoder (E), property decoder (P), and image generator (G) are not depicted to avoid congested illustration.

image I_s collected from a very sparse property space. In contrast, by sampling the target property α_t from a continuous property space, the input image I_s can be tuned to any α_t . In addition, we train with the paired data of generated image \hat{I}_s and source image I_s by checking whether $\eta(\hat{I}_s) = \eta(I_s)$. As illustrated in Figure 2, the generated sample I_t has its corresponding self-generated label α_t . The self-generated data ensures any feasible property values can be explored during training by randomly sampling α_t , so the model can be freed from relying on a large amount of labeled training samples and makes it possible to disentangle the tuning property with the invariant property. It also allows the model to train with diverse and arbitrarily sparsely distributed data. For example, starting from one single face image of a person, the self-generated data can leverage the transferred knowledge from another person with a different face pose to generate self-supervised data for self-learning.

3.3. Checking with Capsule Representation

To disentangle the tunable property vs invariant ones, we need to have a self-attention mechanism which can select the relevant information from each image. For example, as the face turns, a self-attention mechanism is needed to relocate the key facial features for checking. Furthermore, the invariant properties could be diverse with multimodal distributions. We adopt the capsule network to support the

disentanglement. We use the capsule representation for two purposes: 1) dynamic routing of capsule finds where is the relevant spatial information, so we can be adaptive to the image variation due to tuning; 2) by using the vector representation, the capsule can encode a multimodal distance by condensing the error measure to the length of the vector.

Invariant Check. As shown in Figure 2, we adopt a capsule network to project the invariant information extracted from an image to a low-dimensional vector output. A crucial design of our system is using capsule to guide the image generation process, as well as checking the invariant properties. For the example of head rotation, we aim to make sure that the face identity is maintained during the tuning process. The invariant information β is modeled by the capsule network as a vector representation. We establish the loss function for the invariant check with capsule vectors by:

$$\mathcal{L}_i = \|C(E(I_t))\|_2 - \|C(E(I_s))\|_2 \quad (2)$$

where $C(E(I_s))$ is the capsule vector of I_s , $C(E(I_t))$ is the capsule vector of I_t and $\|\cdot\|_2$ indicates L2-norm. During back propagation, the gradient from the capsule updates the network to make sure that the two images I_s and I_t share the same invariant information, as shown in Figure 2.

Tunable property Check. As shown in Figure 3, a capsule network is also proposed for tunable property check. A arbitrary image (e.g. I_t) is first fed into a backbone network, and then the output feature maps are processed by a capsule

324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344

378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431

432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
network based estimator (*e.g.* E_1 in the Figure) to extract the tunable property related information and turn it into a capsule vector whose L2-norm serves as the tunable property state estimation. Comparing with directly regressing from the feature map, the capsule network selects most relevant information that can suppress irrelevant information such as lighting and shading. As explained in the following subsection, we choose to use a multi-branch structure for the property estimator in our TunerGAN, thus each branch η_i contains a capsule-based estimator which generates an estimated property $\hat{\alpha}_t^i$ that can be used for computing the following loss:

$$\mathcal{L}_p^i = \|\eta_i(I_t) - \alpha_t\|_1 = \|\hat{\alpha}_t^i - \alpha_t\|_1. \quad (3)$$

3.4. Multi-Branch Estimator and Discriminator

To make sure we are generating a real image from the desired distribution, additional adversarial loss, and tunable property checking loss are adopted by introducing a Multi-branch Estimator and Discriminator module (MBED). It has two functionalities: 1) property checking and 2) image quality checking, realized by estimators and discriminators, respectively. To prevent overfitting, we deploy partial parameter sharing between “property checking” and “image quality checking” to facilitate gradient propagation between them. The GAN training often suffers from vanishing gradient and mode collapse. To stabilize the training, we design an ensemble of multiple branches for both estimator and discriminator. We observe in our experiments, as also shown in Figure 3, the multi-branch structure can force the proper checking to be more precise and the discriminators to focus on different facial features. Each estimator use a capsule network for property checking with \mathcal{L}_p^i (Eq. 3) and the overall property checking loss is just a simple sum of the losses from the branches: $\mathcal{L}_p = \sum_{i=1}^N \mathcal{L}_p^i$, where N is the number of branches and it is set to 4 in our experiments. For image quality checking, we use the PatchGAN loss [13] to constrain each discriminator.

4. Experiments

4.1. Datasets

UTKFace [37] is a large-scale face dataset with long age span (range from 0 to 116 years old). The dataset consists of over 20,000 face images with annotations of age, gender, and ethnicity. It is a challenging task to generate continuous facial images with age changes.

RaFD [15] contains pictures with 67 identities displaying 8 emotional expressions. In our experiments, we use identities between 1 and 60 identities for training and the rest for evaluation. Note we only use the frontal faces to capture the emotion distribution.

FEI [29] is a Brazilian face database. There are 14 images for each of 200 individuals, resulting in a total of 2800 im-

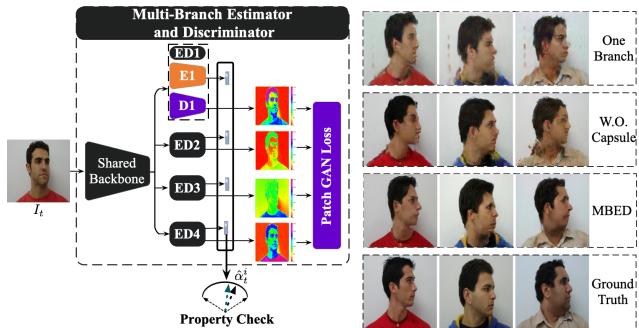


Figure 3. On the left, we show the structure of our multi-branch Estimator and Discriminator (MBED). On the right, we compare the performance of the property tuning network trained by MBED, one capsule with one discriminator, and the multi-discriminator without capsule. All images are generated by the same head pose as the ground truth. The MBE can teach the generator to give better images even for a big head pose changing.

ages. The profile rotation is up to about 180 degrees with 10 discrete samples. The first 160 identities are used for training and others for evaluation.

Foggy Kitty [7] provides 7 categories fog images with different visibility: 30m, 40m, 50m, 75m, 150m, 375m, 750m. We follow the original train/test split.

CelebA [19] is a large face dataset containing $\sim 210,000$ images with 40 different attribute annotations. We perform multiple-property tuning on this dataset.

4.2. Evaluation Metric and Experimental Setup

Fréchet Inception Distance (FID) [8] is proposed to compute the similarity between the generated sample distribution and real data distribution. A lower FID score indicates better-generated image quality.

Learned Perceptual Image Patch Similarity (LPIPS) is proposed by [36], which computes the perceptual similarity between two image patches. Lower LPIPS suggests two image patches have closer perceptual similarities.

MSE, **RMSE** [31], **PSNR** [11] and **SSIM** [11] are used to measure the distance for the paired data. The lower MSE and RMSE values indicate better performance while the higher PSNR and SSIM indicate a closer similarity.

Classification Acc [4] measures the distance on the feature expression. We use the off-the-shelf classification model with Inception-V3 architecture to compute the accuracy.

To evaluate our TunerGAN, we compare with state-of-the-art methods in different tasks. For paired samples, we compare our method with paired methods: Pix2pix [13] and PG² [23]. As for unpaired methods, we compare our method with StarGAN [4], STGAN [18], and StarGAN+MUNIT.

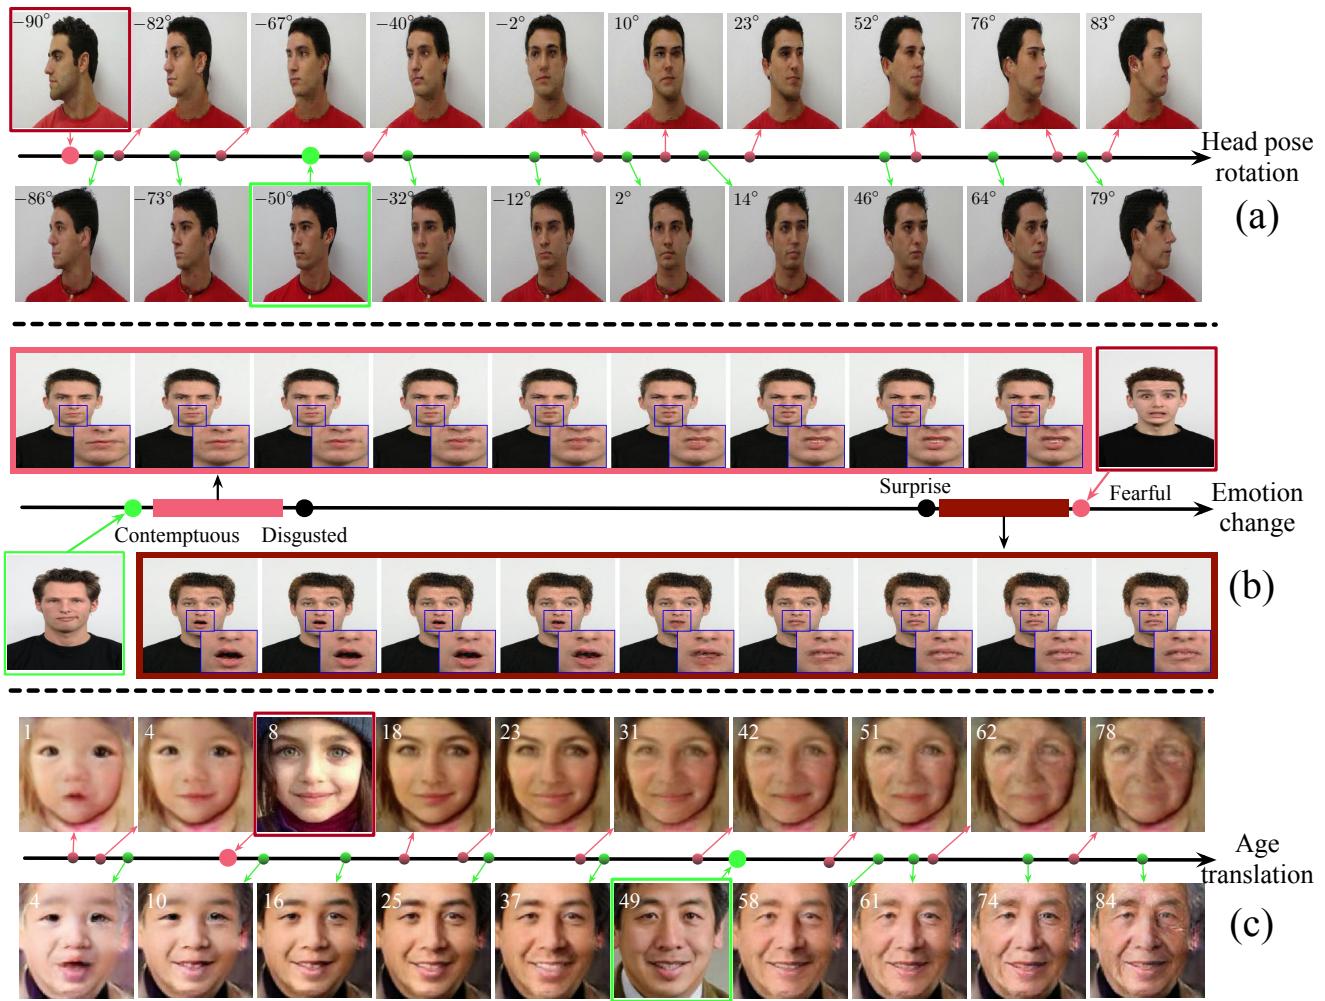


Figure 4. The visual translation results of TunerGAN on FEI, RaFD, and UTKface datasets. The images in red and green boxes show the real input samples, and other images are generated samples. The upper left corner of the image is marked with the pose angle of the generated images. The carnations and green dots represent the projection in the latent space of the generated samples.

4.3. Single-property tuning

We first evaluate our method on the continuous **head pose rotation** task using the FEI dataset. Figure 4(a) exhibits the results of the tuning process of the head pose. The images in red and green boxes are real input images while others are generated samples. We project the images into head pose space and use a red dot for the first image and green for the second. The results demonstrate that our method successfully generates images with an arbitrary head pose, which indicates our property checking can guide the property tuning network to generate data with target property. At the same time, the image with the target pose can still retain the same identity of the input image, which shows our invariant checking can anchor the invariant information during image manipulation. To perform a quantitative comparison with other methods, we generate 10 images by head pose with ground truth images. The

Pix2pix [13] is compared with by using an additional one-hot encoding to generate images. We conduct PG² by using the facial landmarks extracted by dlib [14] as pose guidance. StarGAN and STGAN are compared with for the generation of discrete samples. Table 1 reports the quantitative comparison of the above methods. We outperform other methods on some of the evaluation metrics, no matter whether they are trained with paired or unpaired data. The visualization of the comparison results is shown in Figure 5. Compared with StarGAN and STGAN, the image generated by TunerGAN can retain the same emotion, identity as the input, and have more real details and fewer artifacts on the facial features.

In addition, we conduct the continuous **emotional expression tuning** on RaFD [15], which is similar to emotion changing as described in GANimation [26], without requiring a lot of intermediate samples and hand-

648 Table 1. Quantitative comparison of viewpoint transfer (facial pose translation) on FEI dataset. The symbol \uparrow (\downarrow) indicates that the larger
 649 (smaller) the value, the better the performance. \dagger represents the methods trained with paired samples. 702
 650 703

Method	SSIM \uparrow	PSNR \uparrow	MSE \downarrow	RMSE \downarrow	LPIPS \downarrow	FID \downarrow	Classification Acc \uparrow
Pix2pix \dagger [13]	0.9215	16.90	0.0162	0.1232	0.2273	74.31	0.9450
PG 2† [23]	0.9354	17.34	0.0145	0.1160	0.1887	70.70	0.9525
StarGAN [4]	0.8812	15.17	0.0247	0.1512	0.2567	93.26	0.9325
STGAN [18]	0.8931	15.44	0.0232	0.1483	0.2514	76.56	0.9350
TunerGAN	0.9219	16.68	0.0199	0.1288	0.2123	66.56	0.9650



668 Figure 5. Visual result examples of supervised and unsupervised
 669 methods on FEI dataset. 709

671 crafted labels. We tune the continuous expression among
 672 8 sparse emotional samples, following the guidelines of
 673 the FACS system [5] and the emotional changing procedure:
 674 Neutral \rightarrow Contemptuous \rightarrow Disgusted \rightarrow Happy \rightarrow Angry \rightarrow Sad
 675 \rightarrow Surprise \rightarrow Fear. Figure 4(b) shows continuous emotional
 676 changing procedure between two different emotional
 677 representations. The quality of the tunable property of
 678 emotion can be checked especially from the visualization
 679 of the pose of the mouth. From the illustrations, we can
 680 see that the property is continuously changed along with
 681 the emotion space, while identity remains invariant. We
 682 provide the visual and quantitative comparisons of all the
 683 methods in the supplementary file. 720

684 Finally, **age manipulation** aims to generate the past and
 685 future images given only a single image of a certain person.
 686 The generated image should have the same identity
 687 and gender but with different ages. The results of tuning
 688 the age property are shown in Figure 4(c). The images in
 689 red and green boxes are real samples and others are
 690 generated samples by tuning age on them. Our method can
 691 realize a continuous age changing process with one image,
 692 even if there is only one image for some persons in training.
 693 To evaluate the image manipulation performance we
 694 use the online face detection API ¹ to estimate the age of
 695 the generated images. Table. 2 shows the quantitative
 696 comparison of different methods. For evaluation, we compute
 697 the average age error (AE) between the estimated age and
 698 the desired age label, and the face similarity (FS) between
 699 the generated image and the real image to evaluate whether
 700

701 ¹<https://www.faceplusplus.com.cn/face-detection>

702 the identity is retained. We generate 40 images (age ranging
 703 between 20 and 59) for each image. We evaluated 100
 704 images from different persons, and the labels of input images
 705 are also between 20 and 59. We divide the input images into four ranges 20-29, 30-39, 40-49, and 50-59. As
 706 shown in Table. 2, we achieve the lowest average age error
 707 and FID compared with StartGAN and STGAN on all of the
 708 age ranges at the cost of a few percent relative loss in FS.
 709 StarGAN and STGAN just maintain the similar appearance
 710 of the input images so that the face similarity is high. 711

712 Besides tuning face attributes, we also conduct experiments
 713 on scene data. We show an example of applying
 714 TunerGAN to fog thickness tuning on the Foggy Kitti
 715 dataset [7], with visual results shown in Figure 6. More
 716 visual generation results and comparison with other methods
 717 can be found in our supplementary file. 718



719 Figure 6. Visual results on fog thickness tuning over an input image
 720 with 375m visibility. The image in red box is the input. 721

722 4.4. Multiple-property tuning 738

723 Our algorithm can also be modified to handle multiple- 739
 724 property tuning by adding parallel input/output tuning 740
 725 messages, one per property. We conduct experiments on the 741
 726 CelebA dataset [20]. We select 6 parallel properties for tuning: 742
 727 neutral to smile, no beard to goatee, bald to bangs, 743
 728 black hair to brown hair, young to old, and male to female. 744
 729 Figure 7 shows the visual generation results. The first 6 745
 730 rows show the tuning the six parallel face attributes 746
 731 individually, the 7th and 8th rows represent the results of tuning 747
 732 two properties at the same time. We tune all the properties 748
 733 at the same time in the last row. 749

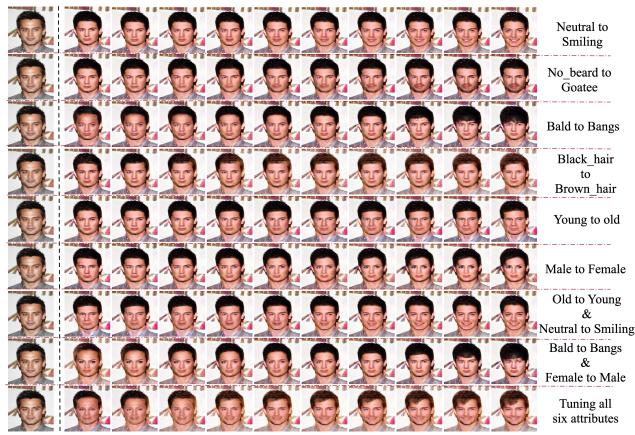
750 We conduct ablation experiments on facial head pose 751
 752 rotation task, with results shown in Table. 3 to analyze 753
 753 the effectiveness of the proposed modules. **Multi-Branch**
 754 **Estimator:** we evaluate it with four different settings:
 1. one-branch capsule regression and discriminator (w/o

756 Table 2. Quantitative comparison of viewpoint transfer (facial pose translation) on UTKface dataset. **AE** and **FS** represent the average age
 757 estimation error and face similarity, separately. 810
 758 811

Method	FID \downarrow	20-29(AE \downarrow / FS \uparrow)	30-39(AE \downarrow / FS \uparrow)	40-49(AE \downarrow / FS \uparrow)	50-59(AE \downarrow / FS \uparrow)
StarGAN [4]	70.26	4.129/0.9782	4.198/0.9672	4.569/0.9718	4.657/0.9741
STGAN [18]	62.36	3.387/0.9819	2.872/0.9796	3.873/0.9862	3.671/0.9827
TunerGAN	50.80	2.346/0.9632	1.984/0.9736	2.468/0.9582	2.580/0.9481

763 Table 3. Quantitative comparison of different settings on FEI dataset. 817

ID	Method	SSIM \uparrow	PSNR \uparrow	MSE \downarrow	RMSE \downarrow	LPIPS \downarrow	FID \downarrow	Classification Acc \uparrow
Exp1	w/o MB	0.8813	15.11	0.0273	0.1652	0.2612	115.6	0.9025
Exp2	w/ MB(N=2)	0.9126	15.97	0.0226	0.1401	0.2291	76.83	0.9325
Exp3	w/o capsule	0.9089	15.79	0.0239	0.1405	0.2285	82.63	0.9425
Exp4	StarGAN+MUNIT	0.9015	15.27	0.0236	0.1501	0.2418	78.22	0.9375
Exp5	TunerGAN	0.9219	16.68	0.0199	0.1288	0.2123	66.56	0.9650



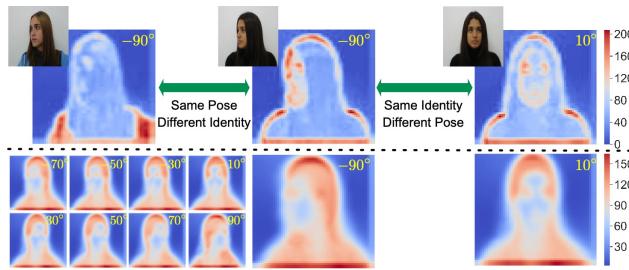
771 Figure 7. Tuning multiple properties on CelebA. 824

772 MB), 2. two-branch capsule regression and discriminator
 773 (w/MB(N=2)), 3. four-branch regression without capsule
 774 and discriminator and 4. four branches as proposed. They
 775 correspond to Exp1, Exp2, Exp3 and Exp5, respectively.
 776 We can observe the multi-branch architecture can promote
 777 the performance of both discrimination and regression.

778 **Capsule Regression:** Comparing Exp3 with Exp5, we can
 779 see that adding the capsule regression to check the property
 780 improves the image generation performance on all metrics.
 781 This demonstrates that the disentanglement of representation
 782 can help the generator to capture more informative details.
 783 Exp1 and Exp5 show that we obtain more accurate
 784 regression with a multi-branch structure.

802 4.5. Ablation study 803

804 **Invariant feature:** To illustrate that our invariant feature
 805 does not end up learning a constant vector, we have visu-
 806 alized the activation of the instance-invariant capsule. We
 807 compare two settings: 1) pose invariant identity change,
 808 2) identity invariant pose change. The visual results are shown
 809 in Figure 8. From the visualization, we observe that the in-

810 Figure 8. Upper: Learned invariant features for 3 exemplar images.
 811 Lower: Mean feature map per pose angle of FEI dataset. 812

812 variant representation learned in these two tasks are very
 813 different. For pose change, it focuses on facial features
 814 (nose, eyes), while for identity change, it focuses on hair
 815 and cloth. We also computed the mean feature map per pose
 816 angle to visualize the learned invariance. 817

817 To illustrate that our method is not just a combination of
 818 StarGAN and MUNIT by applying the content-preserving
 819 loss to StarGAN, we implemented the StarGAN+MUNIT
 820 in Exp4 (see Table 3). TunerGAN outperforms the Star-
 821 GAN+MUNIT by automatically sampling data-label pairs,
 822 going beyond existing labels. Our coupled capsule network
 823 also provides a stronger property invariant protection than
 824 the content-preserving loss in MUNIT. 825

826 5. Conclusion 827

828 We have proposed an instance-invariant property tun-
 829 able image manipulation method that is trained with
 830 synthetically-generated images. We disentangled the image
 831 representation into the invariant information and property
 832 information. We ensured the instance-invariant informa-
 833 tion to be consistent during the manipulation and checked
 834 whether the image having the target property by two types
 835 of capsule networks. We have demonstrated the proposed
 836 TunerGAN method outperforms all existing algorithms on
 837 instance level property tuning image manipulation tasks. 838

864

References

865

866

867

868

869

870

871

872

873

874

875

876

877

878

879

880

881

882

883

884

885

886

887

888

889

890

891

892

893

894

895

896

897

898

899

900

901

902

903

904

905

906

907

908

909

910

911

912

913

914

915

916

917

- [1] Xi Chen, Yan Duan, Rein Houthooft, John Schulman, Ilya Sutskever, and Pieter Abbeel. InfoGAN: Interpretable representation learning by information maximizing generative adversarial nets. In *NIPS*, pages 2172–2180, 2016. 1
- [2] Ying-Cong Chen, Xiaogang Xu, Zhuotao Tian, and Jiaya Jia. Homomorphic latent space interpolation for unpaired image-to-image translation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2408–2416, 2019. 3
- [3] Zhimin Chen and Yuguang Tong. Face super-resolution through Wasserstein GANs. *arXiv preprint arXiv:1705.02438*, 2017. 3
- [4] Yunjey Choi, Minje Choi, Munyoung Kim, Jung-Woo Ha, Sunghun Kim, and Jaegul Choo. StarGAN: Unified generative adversarial networks for multi-domain image-to-image translation. In *CVPR*, pages 8789–8797, 2018. 1, 3, 5, 7, 8
- [5] Rosenberg Ekman. *What the face reveals: Basic and applied studies of spontaneous expression using the Facial Action Coding System (FACS)*. Oxford University Press, USA, 1997. 6
- [6] Leon A Gatys, Alexander S Ecker, and Matthias Bethge. A neural algorithm of artistic style. *arXiv preprint arXiv:1508.06576*, 2015. 3
- [7] Shirsendu Sukanta Halder, Jean-François Lalonde, and Raoul de Charette. Physics-based rendering for improving robustness to rain. In *IEEE/CVF International Conference on Computer Vision*, 2019. 5, 7
- [8] Martin Heusel, Hubert Ramsauer, Thomas Unterthiner, Bernhard Nessler, and Sepp Hochreiter. GANs trained by a two time-scale update rule converge to a local Nash equilibrium. In *NIPSs*, pages 6626–6637, 2017. 5
- [9] Geoffrey E Hinton, Alex Krizhevsky, and Sida D Wang. Transforming auto-encoders. In *International Conference on Artificial Neural Networks*, pages 44–51, 2011. 3
- [10] Geoffrey E Hinton, Sara Sabour, and Nicholas Frosst. Matrix capsules with em routing. 2018. 3
- [11] Alain Hore and Djemel Ziou. Image quality metrics: Psnr vs. ssim. In *International Conference on Pattern Recognition*, pages 2366–2369. IEEE, 2010. 5
- [12] Xun Huang, Ming-Yu Liu, Serge Belongie, and Jan Kautz. Multimodal unsupervised image-to-image translation. In *European Conference on Computer Vision*, pages 172–189, 2018. 1, 3
- [13] Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, and Alexei A Efros. Image-to-image translation with conditional adversarial networks. In *CVPR*, pages 5967–5976, 2017. 3, 5, 6, 7
- [14] Davis E King. Dlib-ml: A machine learning toolkit. *Journal of Machine Learning Research*, 10(Jul):1755–1758, 2009. 6
- [15] Oliver Langner, Ron Dotsch, Gijsbert Bijlstra, Daniel HJ Wigboldus, Skyler T Hawk, and AD Van Knippenberg. Presentation and validation of the Radboud Faces Database. *Cognition and Emotion*, 24(8):1377–1388, 2010. 5, 6
- [16] Christian Ledig, Lucas Theis, Ferenc Huszár, Jose Caballero, Andrew Cunningham, Alejandro Acosta, Andrew Aitken,

- Alykhan Tejani, Johannes Totz, Zehan Wang, and Wenzhe Shi. Photo-realistic single image super-resolution using a generative adversarial network. In *CVPR*, pages 4681–4690, 2017. 3
- [17] Hsin-Ying Lee, Hung-Yu Tseng, Jia-Bin Huang, Maneesh Singh, and Ming-Hsuan Yang. Diverse image-to-image translation via disentangled representations. In *European Conference on Computer Vision*, pages 35–51, 2018. 3
- [18] Ming Liu, Yukang Ding, Min Xia, Xiao Liu, Errui Ding, Wangmeng Zuo, and Shilei Wen. Stgan: A unified selective transfer network for arbitrary image attribute editing. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3673–3682, 2019. 1, 3, 5, 7, 8
- [19] Ziwei Liu, Ping Luo, Xiaogang Wang, and Xiaou Tang. Deep learning face attributes in the wild. In *Proceedings of International Conference on Computer Vision (ICCV)*, December 2015. 5
- [20] Ziwei Liu, Ping Luo, Xiaogang Wang, and Xiaou Tang. Deep learning face attributes in the wild. In *International Conference on Computer Vision*, pages 3730–3738, 2015. 7
- [21] Jonathan Long, Evan Shelhamer, and Trevor Darrell. Fully convolutional networks for semantic segmentation. In *CVPR*, pages 3431–3440, 2015. 2
- [22] Yu Luo, Yong Xu, and Hui Ji. Removing rain from a single image via discriminative sparse coding. In *ICCV*, pages 3397–3405, 2015. 2
- [23] Liqian Ma, Xu Jia, Qianru Sun, Bernt Schiele, Tinne Tuytelaars, and Luc Van Gool. Pose guided person image generation. In *Advances in Neural Information Processing Systems*, pages 406–416, 2017. 5, 7
- [24] Mehdi Mirza and Simon Osindero. Conditional generative adversarial nets. *arXiv preprint arXiv:1411.1784*, 2014. 3
- [25] Deepak Pathak, Philipp Krahenbuhl, Jeff Donahue, Trevor Darrell, and Alexei A Efros. Context encoders: Feature learning by inpainting. In *CVPR*, pages 2536–2544, 2016. 3
- [26] Albert Pumarola, Antonio Agudo, Aleix M Martinez, Alberto Sanfeliu, and Francesc Moreno-Noguer. Ganimation: Anatomically-aware facial animation from a single image. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 818–833, 2018. 1, 6
- [27] Sara Sabour, Nicholas Frosst, and Geoffrey E Hinton. Dynamic routing between capsules. In *Advances in Neural Information Processing Systems*, pages 3859–3869, 2017. 3
- [28] Casper Kaae Sønderby, Jose Caballero, Lucas Theis, Wenzhe Shi, and Ferenc Huszár. Amortised map inference for image super-resolution. In *ICLR*, pages 1–17, 2017. 3
- [29] Carlos Eduardo Thomaz and Gilson Antonio Giraldi. A new ranking method for principal components analysis and its application to face image analysis. *Image and Vision Computing*, 28(6):902–913, 2010. 5
- [30] Ting-Chun Wang, Ming-Yu Liu, Jun-Yan Zhu, Andrew Tao, Jan Kautz, and Bryan Catanzaro. High-resolution image synthesis and semantic manipulation with conditional gans. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 8798–8807, 2018. 3

918

919

920

921

922

923

924

925

926

927

928

929

930

931

932

933

934

935

936

937

938

939

940

941

942

943

944

945

946

947

948

949

950

951

952

953

954

955

956

957

958

959

960

961

962

963

964

965

966

967

968

969

970

971

- 972 [31] Cort J Willmott and Kenji Matsuura. Advantages of the mean 1026
973 absolute error (mae) over the root mean square error (rmse) 1027
974 in assessing average model performance. *Climate research*, 1028
975 30(1):79–82, 2005. 5 1029
- 976 [32] Chao Yang, Xin Lu, Zhe Lin, Eli Shechtman, Oliver Wang, 1030
977 and Hao Li. High-resolution image inpainting using multi- 1031
978 scale neural patch synthesis. In *CVPR*, pages 6721–6729, 1032
979 2017. 3 1033
- 980 [33] Raymond A Yeh, Chen Chen, Teck Yian Lim, Alexander G 1034
981 Schwing, Mark Hasegawa-Johnson, and Minh N Do. Semantic 1035
982 image inpainting with deep generative models. In *CVPR*, 1036
983 pages 5485–5493, 2017. 2 1037
- 984 [34] Zili Yi, Hao Zhang, Ping Tan, and Minglun Gong. Dual- 1038
985 GAN: Unsupervised dual learning for image-to-image trans- 1039
986 lation. In *International Conference on Computer Vision*, 1040
987 pages 2849–2857, 2017. 3 1041
- 988 [35] He Zhang, Vishwanath Sindagi, and Vishal M Patel. Image 1042
989 de-raining using a conditional generative adversarial net- 1043
990 work. *arXiv preprint arXiv:1701.05957*, 2017. 2 1044
- 991 [36] Richard Zhang, Phillip Isola, Alexei A Efros, Eli Shechtman, 1045
992 and Oliver Wang. The unreasonable effectiveness of deep 1046
993 features as a perceptual metric. In *CVPR*, pages 586–595, 1047
994 2018. 5 1048
- 995 [37] Zhifei Zhang, Yang Song, and Hairong Qi. Age progres- 1049
996 sion/regression by conditional adversarial autoencoder. In 1050
997 *CVPR*, pages 5810–5818, 2017. 5 1051
- 998 [38] Jun-Yan Zhu, Taesung Park, Phillip Isola, and Alexei A 1052
999 Efros. Unpaired image-to-image translation using cycle- 1053
1000 consistent adversarial networks. In *International Conference 1054
1001 on Computer Vision*, pages 2223–2232, 2017. 1, 3 1055
1002 1056
1003 1057
1004 1058
1005 1059
1006 1060
1007 1061
1008 1062
1009 1063
1010 1064
1011 1065
1012 1066
1013 1067
1014 1068
1015 1069
1016 1070
1017 1071
1018 1072
1019 1073
1020 1074
1021 1075
1022 1076
1023 1077
1024 1078
1025 1079